

# Lecture 6: Eligibility Traces

Max(Chong) Li

# Outline

---

- TD( $\lambda$ ) Learning
- TD( $\lambda$ ) Control: Sarsa( $\lambda$ ) & Q( $\lambda$ )
- The Unified View of RL Solutions

\*Some materials are modified from David Silver's RL lecture notes


# Outline

---

- TD( $\lambda$ ) Learning
- TD( $\lambda$ ) Control: Sarsa( $\lambda$ ) & Q( $\lambda$ )
- The Unified View of RL Solutions

# TD(0) Backup (Refresher)

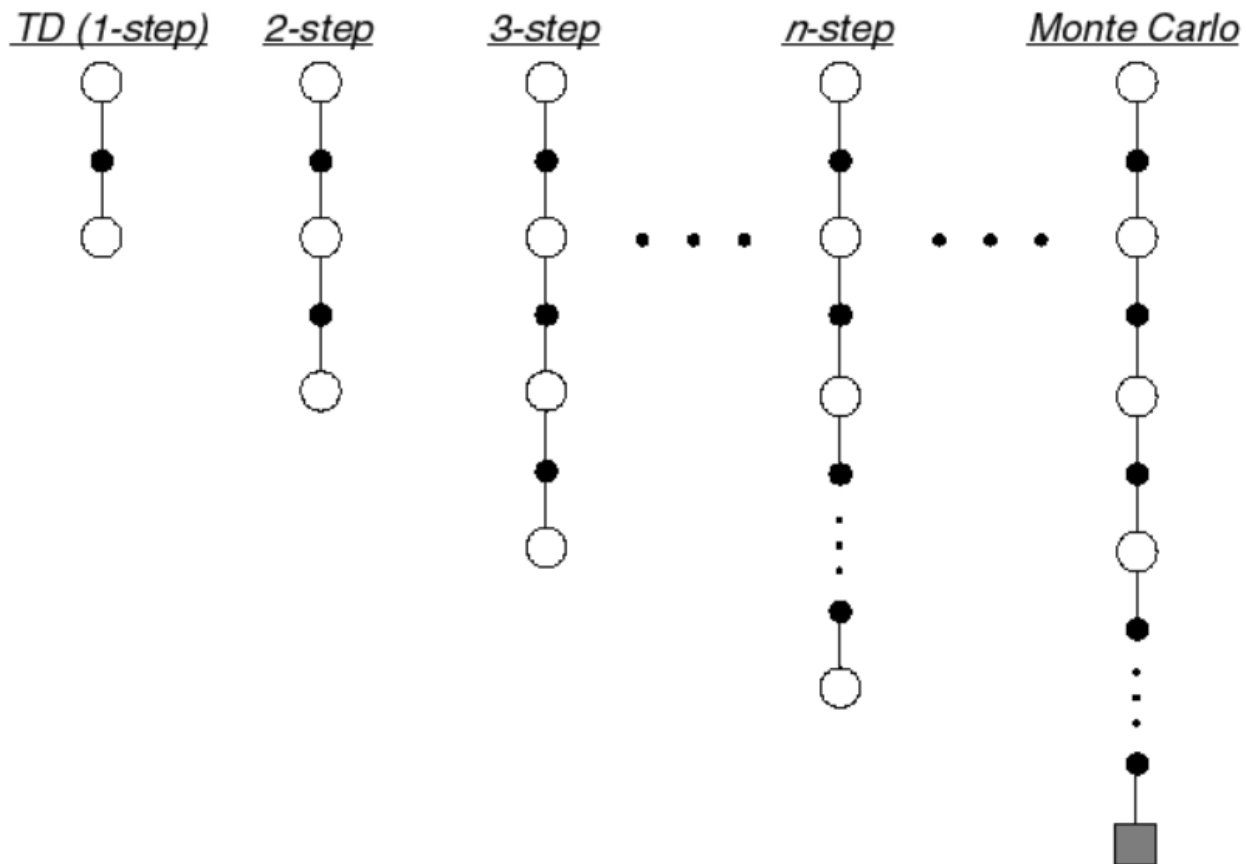
$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$



# n-step prediction

---

Let TD target look  $n$  steps into the future



# n-step return

---

- Consider the following  $n$ -step returns for  $n = 1, 2, \infty$ :

$$\begin{array}{ll} n = 1 & (TD) \quad G_t^{(1)} = R_{t+1} + \gamma V(S_{t+1}) \\ n = 2 & G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V(S_{t+2}) \\ \vdots & \vdots \\ n = \infty & (MC) \quad G_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T \end{array}$$

- Define the  $n$ -step return

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$

- $n$ -step temporal-difference learning


$$V(S_t) \leftarrow V(S_t) + \alpha \left( G_t^{(n)} - V(S_t) \right)$$

# Averaged n-step return

- We can average  $n$ -step returns over different  $n$   
e.g. average the 2-step and 4-step returns

$$\frac{1}{2}G^{(2)} + \frac{1}{2}G^{(4)}$$

- Combines information from two different time-steps
- Can we efficiently combine information from all time-steps?




# $\lambda$ -return

- The  $\lambda$ -return  $G_t^\lambda$  combines all  $n$ -step returns  $G_t^{(n)}$
- Using weight  $(1 - \lambda)\lambda^{n-1}$

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$


- Forward-view  $\text{TD}(\lambda)$

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$






# Weighting Function



$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

# Forward View TD( $\lambda$ )



- Update value function towards the  $\lambda$ -return
- Forward-view looks into the future to compute  $G_t^\lambda$
- Like MC, can only be computed from complete episodes

# Forward View & Backward View

---

- Forward view provides theory
- Backward view provides mechanism
- Update online, every step, from incomplete sequences


# Eligibility Trace

---

- **Frequency heuristic**: assign credit to most frequent states
- **Recency heuristic**: assign credit to most recent states
- *Eligibility traces* combine both heuristics

$$E_0(s) = 0$$


$$E_t(s) = \gamma\lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$



# Backward View TD( $\lambda$ )

- Keep an eligibility trace for every state  $s$
- Update value  $V(s)$  for every state  $s$
- In proportion to TD-error  $\delta_t$  and eligibility trace  $E_t(s)$

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$
$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$



## TD( $\lambda$ ) and TD(0)

---

- When  $\lambda = 0$ , only current state is updated

$$E_t(s) = \mathbf{1}(S_t = s)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$

- This is exactly equivalent to TD(0) update

$$V(S_t) \leftarrow V(S_t) + \alpha \delta_t$$

# Offline Equivalence of Forward and Backward Views

---

- Updates are accumulated within episode
- but applied in batch at the end of episode

*The sum of offline updates is identical for forward-view and backward-view  $TD(\lambda)$*

$$\sum_{t=1}^T \alpha \delta_t E_t(s) = \sum_{t=1}^T \alpha \left( G_t^\lambda - V(S_t) \right) \mathbf{1}(S_t = s)$$

# On-line TD( $\lambda$ )

---

```
Initialize  $V(s)$  arbitrarily
Repeat (for each episode):
  Initialize  $Z(s) = 0$ , for all  $s \in \mathcal{S}$ 
  Initialize  $S$ 
  Repeat (for each step of episode):
     $A \leftarrow$  action given by  $\pi$  for  $S$ 
    Take action  $A$ , observe reward,  $R$ , and next state,  $S'$ 
     $\delta \leftarrow R + \gamma V(S') - V(S)$ 
     $Z(S) \leftarrow Z(S) + 1$ 
    For all  $s \in \mathcal{S}$ :
       $V(s) \leftarrow V(s) + \alpha \delta Z(s)$ 
       $Z(s) \leftarrow \gamma \lambda Z(s)$ 
     $S \leftarrow S'$ 
  until  $S$  is terminal
```

Notation:  $Z(s) = E(s)$

- No equivalence b/w Backward and Forward.
- Step size is sufficiently small  $\rightarrow$  almost “equivalence”
- Online TD updates are generally better than off-line TD.



# Outline

---

- TD( $\lambda$ ) Learning
- TD( $\lambda$ ) Control: Sarsa( $\lambda$ ) & Q( $\lambda$ )
- The Unified View of RL Solutions

# n-step Sarsa

---

- Consider the following  $n$ -step returns for  $n = 1, 2, \infty$ :

$$n = 1 \quad (\text{Sarsa}) \quad q_t^{(1)} = R_{t+1} + \gamma Q(S_{t+1})$$

$$n = 2 \quad q_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 Q(S_{t+2})$$

$$\vdots$$
$$\vdots$$

$$n = \infty \quad (\text{MC}) \quad q_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$

- Define the  $n$ -step Q-return

$$q_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q(S_{t+n})$$

- $n$ -step Sarsa updates  $Q(s, a)$  towards the  $n$ -step Q-return

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left( q_t^{(n)} - Q(S_t, A_t) \right)$$


# Forward View Sarsa( $\lambda$ )

- The  $q^\lambda$  return combines all  $n$ -step Q-returns  $q_t^{(n)}$
- Using weight  $(1 - \lambda)\lambda^{n-1}$

$$q_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} q_t^{(n)}$$

- Forward-view Sarsa( $\lambda$ )

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left( q_t^\lambda - Q(S_t, A_t) \right)$$



# Backward View Sarsa( $\lambda$ )

---

- Just like TD( $\lambda$ ), we use **eligibility traces** in an online algorithm
- But Sarsa( $\lambda$ ) has one eligibility trace for each state-action pair

$$E_0(s, a) = 0$$

$$E_t(s, a) = \gamma\lambda E_{t-1}(s, a) + \mathbf{1}(S_t = s, A_t = a)$$

- $Q(s, a)$  is updated for every state  $s$  and action  $a$
- In proportion to TD-error  $\delta_t$  and eligibility trace  $E_t(s, a)$

$$\delta_t = R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha \delta_t E_t(s, a)$$

# Sarsa ( $\lambda$ ) Algorithm


---

```
Initialize  $Q(s, a)$  arbitrarily, for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
Repeat (for each episode):
     $Z(s, a) = 0$ , for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
    Initialize  $S, A$ 
    Repeat (for each step of episode):
        Take action  $A$ , observe  $R, S'$ 
        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)
         $\delta \leftarrow R + \gamma Q(S', A') - Q(S, A)$ 
         $Z(S, A) \leftarrow Z(S, A) + 1$ 
        For all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ :
             $Q(s, a) \leftarrow Q(s, a) + \alpha \delta Z(s, a)$ 
             $Z(s, a) \leftarrow \gamma \lambda Z(s, a)$ 
         $S \leftarrow S'; A \leftarrow A'$ 
    until  $S$  is terminal
```


Notation:  $Z(s, a) = E(s, a)$

# Sarsa ( $\lambda$ ) Gridworld Example


Path taken




Action values increased by one-step Sarsa



Action values increased by 10-step Sarsa



Action values increased by Sarsa( $\lambda$ ) with  $\lambda=0.9$

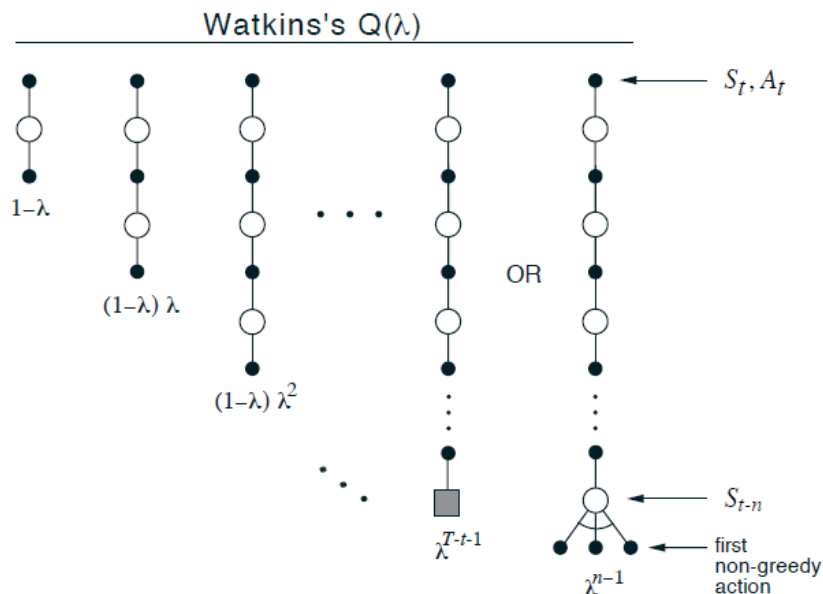


# Q ( $\lambda$ ) Algorithm – Forward View

- We show Watkins's Q( $\lambda$ ). See literature for others Q( $\lambda$ ) algorithms
- n-step Q return: lookahead stops at the first non-greedy action (i.e. exploratory action). If action at  $t+n$  is the first non-greedy action, then the longest backup is toward

$$R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n \max_a Q_t(S_{t+n}, a)$$

- Backup diagram



## Q ( $\lambda$ ) Algorithm – Backward View

---

- One eligibility trace for each state-action pair

$$e_t(s, a) = \begin{cases} 1 + \gamma \lambda e_{t-1}(s, a) & \text{if } s = s_t, a = a_t, Q_{t-1}(s_t, a_t) = \max_a Q_{t-1}(s_t, a) \\ 0 & \text{if } Q_{t-1}(s_t, a_t) \neq \max_a Q_{t-1}(s_t, a) \\ \gamma \lambda e_{t-1}(s, a) & \text{otherwise} \end{cases}$$

- Q update

$$\delta_t = R_{t+1} + \gamma \max_{a'} Q_t(S_{t+1}, a') - Q_t(S_t, A_t)$$

$$Q_{t+1}(s, a) = Q_t(s, a) + \alpha \delta_t E_t(s, a)$$



# Q ( $\lambda$ ) Algorithm

---

```
Initialize  $Q(s, a)$  arbitrarily, for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
Repeat (for each episode):
   $Z(s, a) = 0$ , for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
  Initialize  $S, A$ 
  Repeat (for each step of episode):
    Take action  $A$ , observe  $R, S'$ 
    Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
     $A^* \leftarrow \arg \max_a Q(S', a)$  (if  $A'$  ties for the max, then  $A^* \leftarrow A'$ )
     $\delta \leftarrow R + \gamma Q(S', A^*) - Q(S, A)$ 
     $Z(S, A) \leftarrow Z(S, A) + 1$ 
    For all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ :
       $Q(s, a) \leftarrow Q(s, a) + \alpha \delta Z(s, a)$ 
      If  $A' = A^*$ , then  $Z(s, a) \leftarrow \gamma \lambda Z(s, a)$ 
      else  $Z(s, a) \leftarrow 0$ 
     $S \leftarrow S'; A \leftarrow A'$ 
  until  $S$  is terminal
```

Notation:  $Z(s, a) = E(s, a)$

# Outline

---

- TD( $\lambda$ ) Learning
- TD( $\lambda$ ) Control : Sarsa( $\lambda$ ) & Q( $\lambda$ )
- The Unified View of RL Solutions

# The Unified View of RL Solutions

