

Lecture 8: Policy Gradient

Chonggang Wang

Outline

- Policy Gradient RL
- Actor-Critic Methods
- Policy Gradient w/ Advantage Function

*materials are modified from David Silver's RL lecture notes

Outline

- Policy Gradient RL
- Actor-Critic Methods
- Policy Gradient w/ Advantage Function

Lecture 7 – Value Approximation

- Parameterized value function $\hat{v}(s, \mathbf{w})$ and action-value function $\hat{q}(s, a, \mathbf{w})$ using a parameter column vector \mathbf{w} , to approximate true value functions $V_\pi(s)$ and $Q_\pi(s)$

$$\begin{aligned}\hat{v}(s, \mathbf{w}) &\approx V_\pi(s) \\ \hat{q}(s, a, \mathbf{w}) &\approx Q_\pi(s, a)\end{aligned}$$

- How to obtain $\hat{v}(s, \mathbf{w})$ and $\hat{q}(s, a, \mathbf{w})$
 - Stochastic Gradient Descent: $\Delta \mathbf{w} = \alpha * (V_\pi(s) - \hat{v}(s, \mathbf{w})) * \nabla_{\mathbf{w}} \hat{v}(s, \mathbf{w})$
 - Calculate the gradient $\nabla_{\mathbf{w}} \hat{v}(s, \mathbf{w})$: Represent a state by a feature vector
 - Linear Approach: $\hat{v}(s, \mathbf{w}) = \mathbf{x}(s)^T \mathbf{w}$
 - Non-Linear Approach: Neural Network (e.g., DQN)
- Use the action-value function to determine/select next action based on a pre-configured policy (e.g., greedy, ϵ -greedy)
- Question:** Can we determine next action directly from a parameterized probability function which only depends on the states: **Policy-based RL**

$$\pi_{\theta}(s, a) = p(a|s, \theta)$$

Why Policy-based RL

- Advantages
 - Action probabilities change more smoothly → better convergence
 - Easy to learn stochastic policies → better exploration
 - Avoid “ $\text{argmax}(q(s, a))$ ” → Lower complexity and more effective in high-dimensional or continuous action spaces
- Disadvantages
 - Value-based approaches could be more efficient for a small number of states and actions
 - Hard to get unbiased estimates of policy gradient through sampling
 - Basic policy gradient approaches (e.g., REINFORCE as a Monte Carlo Method) introduce high variance and lead to a lower convergence speed

Policy Objective Functions

- Goal: given policy $\pi_\theta(s, a)$ with parameters θ , find best θ
- But how do we measure the quality of a policy π_θ ?
- In episodic environments we can use the **start value**

$$J_1(\theta) = V^{\pi_\theta}(s_1) = \mathbb{E}_{\pi_\theta} [v_1]$$

- In continuing environments we can use the **average value**

$$J_{avV}(\theta) = \sum_s \underline{d^{\pi_\theta}(s)} V^{\pi_\theta}(s)$$

- Or the **average reward per time-step**

$$J_{avR}(\theta) = \sum_s \underline{d^{\pi_\theta}(s)} \sum_a \pi_\theta(s, a) \mathcal{R}_s^a$$

where $d^{\pi_\theta}(s)$ is **stationary distribution** of Markov chain for π_θ

Policy Optimization

- Policy based reinforcement learning is an **optimisation** problem
- Find θ that maximises $J(\theta)$
- Similar to the value based function approximation, we focus on gradient method
 - Gradient is a key to connect neural network with RL algorithms
- Other approaches are possible



Gradient Descent (recap)

- Let $J(\mathbf{w})$ be a differentiable function of parameter vector \mathbf{w} , a column vector
- Define the *gradient* of $J(\mathbf{w})$ to be

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \begin{pmatrix} \frac{\partial J(\mathbf{w})}{\partial w_1} \\ \vdots \\ \frac{\partial J(\mathbf{w})}{\partial w_n} \end{pmatrix} \quad \begin{array}{l} \text{Partial} \\ \text{Derivatives} \\ \text{with} \\ \text{respect} \\ \text{to } \mathbf{w} \end{array}$$

- To find a local minimum of $J(\mathbf{w})$
- Adjust \mathbf{w} in direction of -ve gradient to reduce $J(\mathbf{w})$ (i.e., the Value Error (VE))

$$\Delta \mathbf{w} = -\frac{1}{2} \alpha \nabla_{\mathbf{w}} J(\mathbf{w})$$



Gradient Ascent

- $J(\theta)$ is the policy-related performance function to be maximized
- $J(\theta)$ is differentiable with respect to θ , which is a column parameter vector

$$\theta = (\theta_1, \dots, \theta_n)^T$$

- $\nabla_{\theta} J(\theta)$: The Gradient of $J(\theta)$ can be calculated

$$\nabla_{\theta} J(\theta) = \begin{pmatrix} \frac{\partial J(\theta)}{\partial \theta_1} \\ \cdot \\ \cdot \\ \frac{\partial J(\theta)}{\partial \theta_n} \end{pmatrix} \quad \begin{array}{l} \text{Partial Derivatives} \\ \text{with respect to } \theta \end{array}$$

- **Gradient Ascent**

- To find the optimal parameter vector which maximizes $J(\theta)$

$$\Delta \theta = \alpha * \nabla_{\theta} J(\theta)$$

Score Function

- We now compute the policy gradient *analytically*
- Assume policy π_θ is differentiable whenever it is non-zero and we know the gradient $\nabla_\theta \pi_\theta(s, a)$
- **Likelihood ratios** exploit the following identity

$$\begin{aligned}\nabla_\theta \pi_\theta(s, a) &= \pi_\theta(s, a) \frac{\nabla_\theta \pi_\theta(s, a)}{\pi_\theta(s, a)} \\ &= \pi_\theta(s, a) \nabla_\theta \underline{\ln \pi_\theta(s, a)}\end{aligned}$$

- The **score function** is $\nabla_\theta \ln \pi_\theta(s, a)$

Example: Softmax Policy

- We will use a softmax policy as a running example
- Weight actions using linear combination of features $\phi(s, a)^T \theta$
- Probability of action is proportional to exponentiated weight

$$\pi_{\theta}(s, a) = p(a|s, \theta) = \frac{e^{\phi(s, a)^T \theta}}{\sum_a e^{\phi(s, a)^T \theta}}$$

- The score function is

$$\underline{\nabla_{\theta} \ln \pi_{\theta}(s, a)} = \phi(s, a) - \mathbb{E}_{\pi_{\theta}} [\phi(s, \cdot)]$$

Example: Softmax Policy

- Proof:

$$\nabla \ln \pi_{\theta}(s, a) = \nabla \ln \frac{e^{\phi(s, a)^T \theta}}{\sum_a e^{\phi(s, a)^T \theta}} \quad \text{--Eq. (1)}$$

$$\nabla \ln \left(\frac{x}{y} \right) = \nabla \ln(x) - \nabla \ln(y) \quad \text{--Eq. (2)}$$

$$\nabla \ln(g(x)) = \frac{\nabla g(x)}{g(x)} \quad \text{--Eq. (3)}$$

$$= \phi(s, a) - \frac{\sum_a (e^{\phi(s, a)^T \theta} \phi(s, a))}{\sum_a e^{\phi(s, a)^T \theta}} \quad \text{--Eq. (4)}$$


$$= \phi(s, a) - \sum_a \frac{e^{\phi(s, a)^T \theta}}{\sum_a e^{\phi(s, a)^T \theta}} \phi(s, a) \quad \text{--Eq. (5)}$$

$$= \phi(s, a) - \mathbb{E}_{\pi_{\theta}}[\phi(s, \cdot)] \quad \text{--Eq. (6)}$$

Example: Gaussian Policy

- In continuous action spaces, a Gaussian policy is natural
- Mean is a linear combination of state features $\mu(s) = \phi(s)^\top \theta$
- Variance may be fixed σ^2 , or can also be parametrised
- Policy is Gaussian, $a \sim \mathcal{N}(\mu(s), \sigma^2)$ $\pi_\theta(a|s, \theta) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(a-\mu(s,\theta))^2}{2\sigma^2}}$
- The score function is

$$\nabla_\theta \ln \pi_\theta(s, a) = \frac{(a - \mu(s))\phi(s)}{\sigma^2}$$



One-Step MDP

- Consider a simple class of **one-step** MDPs
Starting in state $s \sim d(s)$
Terminating after one time-step with reward $r = \mathcal{R}_{s,a}$
- Use likelihood ratios to compute the policy gradient

$$\begin{aligned} J(\theta) &= \mathbb{E}_{\pi_\theta} [r] \\ &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in \mathcal{A}} \pi_\theta(s, a) \mathcal{R}_{s,a} \\ \nabla_\theta J(\theta) &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in \mathcal{A}} \pi_\theta(s, a) \nabla_\theta \ln \pi_\theta(s, a) \mathcal{R}_{s,a} \\ &= \mathbb{E}_{\pi_\theta} [\nabla_\theta \ln \pi_\theta(s, a) r] \end{aligned}$$

Generalized MDP

- The policy gradient theorem generalises the likelihood ratio approach to multi-step MDPs
- Replaces instantaneous reward r with long-term value $\underline{Q^\pi(s, a)}$
- Policy gradient theorem applies to start state objective, average reward and average value objective

*For any differentiable policy $\pi_\theta(s, a)$,
for any of the policy objective functions $J = J_1, J_{avR}$, or $\frac{1}{1-\gamma} J_{avV}$,
the policy gradient is*

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \ln \pi_\theta(s, a) Q^{\pi_\theta}(s, a)]$$

Policy Gradient Theorem

Monte-Carlo Policy Gradient (REINFORCE)

- Update parameters by stochastic gradient ascent
- Using policy gradient theorem
- Using return G_t as an unbiased sample of $Q^{\pi_\theta}(s_t, a_t)$

High Variance

$$\Delta\theta_t = \alpha \nabla_\theta \ln \pi_\theta(s_t, a_t) G_t$$

Policy Gradient

```
1 function REINFORCE
2   Initialise  $\theta$  arbitrarily
3   for each episode  $\{s_1, a_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T\} \sim \pi_\theta$  do
4     for  $t = 1$  to  $T - 1$  do
5        $\theta \leftarrow \theta + \alpha \nabla_\theta \ln \pi_\theta(s_t, a_t) G_t$ 
6     end for
7   end for
8   return  $\theta$ 
9 end function
```


Outline

- Policy Gradient RL
- Actor-Critic Methods
- Policy Gradient w/ Advantage Function

*materials are modified from David Silver's RL lecture notes

Reducing Variance Using a Critic

- Monte-Carlo policy gradient still has high variance
- We use a **critic** to estimate the action-value function,

$$Q_w(s, a) \approx Q^{\pi_\theta}(s, a)$$

- Actor-critic algorithms maintain *two* sets of parameters
 - Critic** Updates action-value function parameters w
 - Actor** Updates policy parameters θ , in direction suggested by critic
- Actor-critic algorithms follow an *approximate* policy gradient

$$\nabla_\theta J(\theta) \approx \mathbb{E}_{\pi_\theta} [\nabla_\theta \ln \pi_\theta(s, a) Q_w(s, a)]$$

$$\Delta\theta = \alpha \nabla_\theta \ln \pi_\theta(s, a) Q_w(s, a)$$

Example: TD(0)-based Actor-Critic Algorithm

- Simple actor-critic algorithm based on action-value critic
- Using linear value fn approx. $Q_w(s, a) = \phi(s, a)^\top w$

Critic Updates w by linear TD(0)

Actor Updates θ by policy gradient

```
1— function QAC
2—   Initialise  $s, \theta$ 
3—   Sample  $a \sim \pi_\theta$ 
4—   for each step do
5—     Sample reward  $r = \mathcal{R}_s^a$ ; sample transition  $s' \sim \mathcal{P}_{s'}^a$ .
6—     Sample action  $a' \sim \pi_\theta(s', a')$ 
7—      $\delta = r + \gamma Q_w(s', a') - Q_w(s, a)$ 
8—      $\theta = \theta + \alpha \nabla_\theta \ln \pi_\theta(s, a) Q_w(s, a)$ 
9—      $w \leftarrow w + \beta \delta \phi(s, a)$   $\leftarrow w = w + \beta * \delta * (\nabla_w Q_w(s, a))$ 
10—     $a \leftarrow a', s \leftarrow s'$ 
11—   end for
12— end function
```

Problem: Bias in Actor-Critic Algorithms

- Approximating the policy gradient introduces bias
- A biased policy gradient may not find the right solution

- Luckily, if we choose value function approximation carefully
- Then we can avoid introducing any bias
i.e. We can still follow the *exact* policy gradient

Compatible Function Approximation Theorem*

If the following two conditions are satisfied:

*Value function approximator is **compatible** to the policy*

$$\nabla_w Q_w(s, a) = \nabla_\theta \ln \pi_\theta(s, a)$$

Value function parameters w minimise the mean-squared error

$$\varepsilon = \mathbb{E}_{\pi_\theta} [(Q^{\pi_\theta}(s, a) - Q_w(s, a))^2] \quad \frac{\partial \varepsilon}{\partial w} = 0$$

Then the policy gradient is exact,

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \ln \pi_\theta(s, a) Q_w(s, a)]$$

*R. Sutton, et al. "Policy gradient methods for reinforcement learning with function approximation", 2000

Outline

- Policy Gradient RL
- Actor-Critic Methods
- Policy Gradient w/ Advantage Function

Reducing Variance Using a Baseline

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \ln \pi_{\theta}(s, a) Q^{\pi_{\theta}}(s, a)] \quad (\text{Without a Baseline})$$

- We subtract a baseline function $B(s)$ from the policy gradient
- This can reduce variance, without changing expectation

$$\begin{aligned} \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \ln \pi_{\theta}(s, a) B(s)] &= \sum_{s \in \mathcal{S}} d^{\pi_{\theta}}(s) \sum_a \nabla_{\theta} \pi_{\theta}(s, a) B(s) \\ \text{\textit{B(s) is action-independent}} \quad &= \sum_{s \in \mathcal{S}} d^{\pi_{\theta}} B(s) \nabla_{\theta} \sum_{a \in \mathcal{A}} \pi_{\theta}(s, a) \\ &= 0 \end{aligned}$$

- A good baseline is the state value function $B(s) = V^{\pi_{\theta}}(s)$
- So we can rewrite the policy gradient using the **advantage function** $A^{\pi_{\theta}}(s, a)$

$$A^{\pi_{\theta}}(s, a) = Q^{\pi_{\theta}}(s, a) - V^{\pi_{\theta}}(s)$$

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \ln \pi_{\theta}(s, a) A^{\pi_{\theta}}(s, a)] \quad (\text{Using a Baseline})$$

"Variance Reduction for Policy Gradient with **Action-Dependent Factorized Baselines**" - ICLR 2018

Estimating the Advantage Function

- The advantage function can significantly reduce variance of policy gradient
- So the critic should really estimate the advantage function
For example, by estimating *both* $V^{\pi_{\theta}}(s)$ and $Q^{\pi_{\theta}}(s, a)$
- Using two function approximators and two parameter vectors,

$$V_v(s) \approx V^{\pi_{\theta}}(s)$$

$$Q_w(s, a) \approx Q^{\pi_{\theta}}(s, a)$$

$$A(s, a) = Q_w(s, a) - V_v(s)$$

- And updating *both* value functions by e.g. TD learning

Estimating the Advantage Function (cont.)

- For the true value function $V^{\pi_\theta}(s)$, the TD error δ^{π_θ}

$$\delta^{\pi_\theta} = r + \gamma V^{\pi_\theta}(s') - V^{\pi_\theta}(s)$$

is an unbiased estimate of the advantage function

$$\begin{aligned}\mathbb{E}_{\pi_\theta} [\delta^{\pi_\theta} | s, a] &= \mathbb{E}_{\pi_\theta} [r + \gamma V^{\pi_\theta}(s') | s, a] - V^{\pi_\theta}(s) \\ &= Q^{\pi_\theta}(s, a) - V^{\pi_\theta}(s) \\ &= A^{\pi_\theta}(s, a)\end{aligned}$$

- So we can use the TD error to compute the policy gradient

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \ln \pi_\theta(s, a) \delta^{\pi_\theta}]$$

- In practice we can use an approximate TD error

$$\delta_v = r + \gamma V_v(s') - V_v(s)$$

- This approach only requires one set of critic parameters v

Critic at Different Time-Scales

- Critic can estimate value function $V_v(s)$ from many targets at different time-scales

For MC, the target is the return v_t

$$\Delta V = \alpha(\mathbf{G}_t - V_v(s))\phi(s)$$


linear approximation

For TD(0), the target is the TD target $r + \gamma V(s')$

$$\Delta V = \alpha(r + \gamma V(s') - V_v(s))\phi(s)$$

For forward-view TD(λ), the target is the λ -return v_t^λ

$$\Delta V = \alpha(\mathbf{G}_t^\lambda - V_v(s))\phi(s)$$

For backward-view TD(λ), we use eligibility traces

$$\delta_t = r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$$

$$e_t = \gamma \lambda e_{t-1} + \phi(s_t)$$

$$\Delta V = \alpha \delta_t e_t$$

Actor at Different Time-Scales

- The policy gradient can also be estimated at many time-scales

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \ln \pi_{\theta}(s, a) A^{\pi_{\theta}}(s, a)]$$

- Monte-Carlo policy gradient uses error from complete return

$$\Delta\theta = \alpha(\mathbf{G}_t - V_v(s_t)) \nabla_{\theta} \ln \pi_{\theta}(s_t, a_t)$$

- Actor-critic policy gradient uses the one-step TD error

$$\Delta\theta = \alpha(\mathbf{r} + \gamma V_v(s_{t+1}) - V_v(s_t)) \nabla_{\theta} \ln \pi_{\theta}(s_t, a_t)$$

Policy Gradient with Eligibility Traces

- Just like forward-view TD(λ), we can mix over time-scales

$$\Delta\theta = \alpha(\textcolor{red}{G}_t^\lambda - V_v(s_t))\nabla_\theta \ln \pi_\theta(s_t, a_t)$$

where $G_t^\lambda - V_v(s_t)$ is a biased estimate of advantage fn

- Like backward-view TD(λ), we can also use eligibility traces
By equivalence with TD(λ), substituting $\phi(s) = \nabla_\theta \ln \pi_\theta(s, a)$

$$\delta = r_{t+1} + \gamma V_v(s_{t+1}) - V_v(s_t)$$

$$e_{t+1} = \lambda e_t + \nabla_\theta \ln \pi_\theta(s, a)$$

$$\Delta\theta = \alpha\delta e_t$$

- This update can be applied online, to incomplete sequences

Summary of Policy Gradient Algorithms

- The **policy gradient** has many equivalent forms

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \ln \pi_{\theta}(s, a) \textcolor{red}{G}_t] && \text{REINFORCE} \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \ln \pi_{\theta}(s, a) \textcolor{red}{Q}^w(s, a)] && \text{Q Actor-Critic} \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \ln \pi_{\theta}(s, a) \textcolor{red}{A}^w(s, a)] && \text{Advantage Actor-Critic} \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \ln \pi_{\theta}(s, a) \textcolor{red}{\delta}] && \text{TD Actor-Critic} \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \ln \pi_{\theta}(s, a) \textcolor{red}{\delta e}] && \text{TD}(\lambda) \text{ Actor-Critic}\end{aligned}$$

- Each leads to a stochastic gradient ascent algorithm.